

Max Ernst Huisman Gutiérrez / CV

Software engineer building production AI systems — from agent tooling to full-stack platforms.



MAXHUISMAN.SPACE

\$ whoami → ai/ml engineer · location → barcelona · status → available now

build passing tests 350+ passing python 3.11 | 3.12 | 3.13 license MIT

4
PYPI PACKAGES

350+
TESTS SHIPPED

88%
BRANCH COVERAGE

25K+
LAUNCH IMPRESSIONS

About

AI changes everything — I build what's missing while it does. Shipped four hot-fix releases in 48h after rewind's launch — third-year engineering student at **La Salle Barcelona** maintaining **four Python packages on PyPI** for **agentic AI workflows**. Together they form a coherent toolkit: **observe** (rewind), **measure** (spent), **secure** (mcpguard), **audit** (debtX). Hot-fixes were triggered by a **real-session smoke test** that caught two Windows-only bugs CI had missed. **mypy strict**, multi-OS CI green, release notes treated like product changelogs. **Open to AI/ML engineering roles in Barcelona — available now.**

Selected projects

rewind — time-travel debugger for AI agent sessions

pypi v0.1.4 license MIT

```
$ pip install rewindx
$ rewind cc setup
✓ Hooks installed in ~/.claude/settings.json
# Time-travel debugging is now active.
```

Python anyio Click Rich SQLite (WAL) content-addressed storage
pytest ruff mypy strict

- ✓ Captures **every prompt, tool call, file edit and cost** into a per-session **SQLite (WAL) + content-addressed blob store**; enables timeline scrub, rollback, undo, privacy-masked Markdown export.
- ✓ **141 tests · 88% branch coverage** · CI green on Python **3.11/3.12/3.13** across Linux · macOS · Windows.
- ✓ **Two production bugs** caught only by a real-session smoke test — Windows **cp1252 stdout crash** and bash-on-Windows backslash escape — both shipped within **24h**.

STRATUM — regulatory intelligence platform for pharma

Loplop.xyz live

Python FastAPI Next.js PostgreSQL Neo4j Weaviate

- ✓ Full-stack AI platform extracting **FDA Complete Response Letters & EMA objections** for compliance teams — processed **28,961 FDA · 418 CRLs · 3,307 objections · 2,296 EMA** through **NLP + graph knowledge representation**.
- ✓ Multi-store data model: **PostgreSQL** (relational), **Neo4j** (graph), **Weaviate** (vector) — multi-dimensional regulatory search across structured, semantic and relational dimensions.

spent — cost tracker for AI agent sessions

pypi v0.5.1 license MIT

Python Rich SQLite JSONL agent hooks

- ✓ **Real-time terminal dashboard** tracking AI session costs with **model detection, per-model breakdown and token-level analysis**.
- ✓ **220+ tests**; launch post reached **25K+ impressions, 100+ engagements** on LinkedIn.

mcpguard OWASP MCP Top 10 security scanner via AST analysis

PyPI · guardmcp · MIT

debtX A–F static-analysis grader for AI-generated code

PyPI · debtX · MIT

Education & Distinctions

La Salle Barcelona — Universitat Ramon Llull

2023 — Present

Bachelor of Engineering in OTIC (ICT & Telecommunications) — Year 3 of 4

1st place — RSME National Math Competition, 2020

Languages

Spanish NATIVE Catalan NATIVE
English C1 Dutch B1

Contact

email max@maxhuisman.space
phone +34 666 459 920
github github.com/loplop-h
linkedin linkedin.com/in/maxernst-huisman
site maxhuisman.space

Stack

code Python (4 yrs primary), TypeScript, SQL
ai / ml LLM APIs, agent frameworks, MCP, RAG, NLP, prompt engineering
backend FastAPI, Next.js, anyio, Click, Rich
data PostgreSQL, SQLite WAL, Neo4j, Weaviate
tooling pytest, ruff, mypy strict, multi-OS CI

Recent ships

MAY 2026 First PR to Anthropic's MCP SDK — stdio fix · in review
APR 2026 Launched rewind v0.1.4 — 4 hot-fix releases in 48h
MAR 2026 Shipped spent v0.5.1 — 25K+ impressions
JAN 2026 Built STRATUM — 28K+ FDA / 2K EMA pipeline

Currently

BUILDING
rewind v0.2 — making AI-agent sessions inspectable like git history

LEARNING
Designing Data-Intensive Applications · use-the-index-luke.com · DeepLearning.AI ML Specialization

Process

- ✓ Building the AI tooling I wish existed when I started
- ✓ mypy strict — types non-negotiable
- ✓ Real-session smoke tests — CI misses the bugs that bite users
- ✓ Release notes as product changelogs — every patch tells a story

Certifications

ML Specialization Stanford / DeepLearning.AI · 2026
GenAI Engineering Professional IBM · 2026
Claude Code in Action Anthropic · 2026
Claude Certified Architect Foundations In Progress